PyGMG: A High Performance, Portable and Expressive Implementation of Multigrid in Python

And

SKYE: Object Tracking with SEJITS

PyGMG:

ASPIRE is a Research Lab at the University of California, Berkeley for computer architecture/high performance computing with joint projects in collaboration with Lawrence Berkeley National Laboratory. https://aspire.eecs.berkeley.edu/about/ I am a student researcher here since December 2014, employed part time by the Lawrence Berkeley Laboratory. I work under the guidance of Armando Fox , UC Berkeley EECS Prof & co-PI ASPIRE Lab, and Kathy Yelick, UCB EECS Prof and LBL Associate Director of Computing Sciences.

This is a poster I presented at the annual ASPIRE lab conference in Santa Cruz, CA, based on my work in the first half of 2015. In this poster, we discuss an implementation of multigrid algorithms that is portable (architecture independent), that can achieve high performance on multicore CPUs and GPUs, and yet is written sequentially such that all parallel code is abstracted away from the developer.

The original implementation of this algorithm, on the other hand, is used widely in the supercomputing community but is verbose and has to be rewritten for different architectures. Our implementation on the other hand, utilizes a code generation platform that allows us to generate code for a pre-specified platform. This allows the maintenance of one easily readable code base but still have the ability to target various architectures.

A follow up presentation is planned for the next ASPIRE conference in Lake Tahoe in early 2016, in which I will discuss improvements to the project. In addition, in a separate project inspired by this work, my colleagues will be presenting their work on a new Domain Specific Language for stencil computations that utilizes code generation, and how it can be used to enhance this software and parallel software in general.

SKYE:

SKYE is a SEJITS (Selective Embedded Just-In-Time Specialization) enhanced program that uses OpenCV and our lab's code generation system (SEJITS) to perform object-tracking in real time. OpenCV includes a rich library of optimized algorithms, and python developers can easily access its resources through the library's python API However, if they want to add any non-OpenCV scripts to their image processing pipeline, there is no way to do it in python without compromising the speed required by real time systems.

An implementation of SEJITS (ctree), our lab's code generation platform, is used to optimize these custom scripts by generating and executing the fast C and OpenMP equivalents of the scripts at run time. We built this tracking algorithm ourselves, and successfully deployed it onto both a Mac OSX computer (using the webcam to capture video, and the computer for the processing), and a Raspberry Pi (where we captured the video feed with a mobile phone and wirelessly transmitted it to the Raspberry Pi for processing), all in real time. Below is a poster presenting of our work that was presented at the 2015 ASPIRE winter retreat in Lake Tahoe (where we also displayed live demos).

PyGMG:
A SEJITS APPLICATION
A High Performance, Portable, and expressive implementation of Multigrid in Python

Shiv Sundram, Nathan Zhang, Chick Markley, Sam Williams, Kathy Yelick, Armando Fox

Problem

Multigrid is numerical technique used to efficiently solve elliptic PDE's (eg Helmholtz, Poisson). Mathematically, this is equivalent to solving a particular sparse matrix equation in $O(N)$ time (where N is the size of the matrix). The existing C implementation (HPGMG) from LBNL is a finite-volume algorithm which is efficient but extremely verbose. Furthermore, it is heavily hardcoded/optimized for a very specific platform (not portable)

Motivation for Python Implementation

3 Goals: **Portability, Efficiency, Readability**
- **Portability**: Should be able to use different platforms (supercomputer, GPU, multicore CPU, or any combination)using one code base (using SEJITS to generate code for each platform)
- **Efficiency**: Should be as fast as HPGMG, and faster to develop/improve
- **Readability**: parallelization of code should be abstracted away from developers and users

Implementation details

- Sequence of stencil computations on a grid
- Stencils/Operators (eg smoothers, interpolators) are modular and adhere to a simple interface.
- Build various data structures wrapped around tuples and numpy arrays to represent grid
- Smoothers: Weighted Jacobi and Chebyshev
- Bottom Solvers: BigCStab, or just repeated smoothing
- Boundary Conditions: Supports both Periodic and Dirichlet

Figure 1: Visualization of stencil computation
Source: http://en.wikipedia.org/wiki/Stencil_code

Algorithm: V-cycle

Multigrid can be decomposed into a series of 3D stencil computations called V-Cycles
- Interpolation/restriction of grid x
- Residual & sparse/symmetric matrix multiply
- Iterative solving of Ax=b (smoothing)

```
Function MGV ( b(i), x(i) )
  //Multigrid V-Cycle
  //Solve T(i)*x(i) = b(i)
  //given b(i) & initial guess for x(i)
  //return improved x(i)
  //let i=log2(size(side(grid))

  if (i = 1)
    //compute exact solution x(1) of P(1) only 1 unknown
    return x(1)
  else
    //smooth/improve solution
    x(i) = Smooth (b(i), x(i))
    r(i) = T(i)*x(i) - b(i) //residual

    r(i-1) = Restrict(r(i))
    x(i-1) = MGV(r(i-1), 0) //recursive call
    d(i)   = Interpolate(x(i-1))

    x(i) = x(i) - d(i) //correct fine grid solution
    x(i) = Smooth(b(i), x(i)) //smooth again
    return x(i)
```
(Algorithm form Jim Demmel's CS267 Lecture Slides)

Figure 2
Cell-centered restriction

"MG V-cycle"
restriction
interpolation

Figure 3
V-cycle

Algorithm: F-cycle

Initial Restriction of RHS

Smooth
Residual
Restrict
Bottom Solve
Interpolate

Figure 4

Can reduce error in solution by at least 1 bit with each call to an F-Cycle, which is simply a sequence of chained V-cycles of increasing initial grid precision

Results: HPGMG vs PyGMG

HPGMG	PyGMG
Hardcoded to only solve 3D grid	Can operate on grids of an arbitrary # of dimensions
MPI and OpenMP scattered throughout code	All code is sequential and thus clearly descriptive of algorithm. No MPI or OpenMP
Can only run code on CPUS & supercomputers	Can theoretically use SEJITs to run code on any computational platform
Equations for generating inputs to grid must be differentiated/manipulated separately and then written into code	Manipulation and differentiation done automatically by Sympy, a computer algebra solver
High Performance	High Performance once coupled with SEJITS

Future Work

Next Steps
- Code generation using SEJITS to achieve Portability
- Bechmark SEJITS generated code against HPGMG to show Efficiency

Other possible areas of interest
- Applying Z-order indexing to grid for increased cache locality
- Overload operators (MATLAB style syntax) instead of using function calls for manipulating grids to achieve even greater Readability

HPGMG

BERKELEY LAB

SKYE | Object Tracking with SEJITS

Mihir Patil | Shiv Sundram | Chick Markley | Armando Fox

WHAT IS SKYE?

SKYE is a SEJITS-enchanced program that uses OpenCV to perform object tracking. While OpenCV has highly-optimized computer vision code, sometimes custom scripts are necessary in object detection; SEJITS is used in SKYE to optimize these scripts.

OBJECT TRACKING

The primary utility of SKYE is as an object tracking system. Originally, the program was a rather basic object-tracking program using only OpenCV function calls. This was fast, but suboptimal, as shown by the blue bounding boxes in the images below.


Figure 1. Consecutive frames of SKYE processing for a video that features a black sock as the sole projectile. In each image, the blue box corresponds to the object detected by the OpenCV-only algorithm, while the red box corresponds to the improved tracking using the "screener" enhancement.

However, after adding a sobel-based "screener" function, that eliminates the residual difference in the image, and decreases the size of the bounding box, shown in red in the images above.

SKYE PROCESSING FLOW

Figure 2. The processing flow for SKYE. The compute intensive "screener" function improves detection performance, and does not introduce a significant slow-down when optimized with SEJITS.

SEJITS PERFORMANCE

After adding in the custom screener code, there was a significant decrease in performance, as shown by the table below.

VERSION	FPS
OpenCV Only	6.99
OpenCV + Custom	0.392
OpenCV + Custom SEJITS	**5.08**

Table 1. The processing speeds, in frames per second of the various versions of development of SKYE.

However, after optimizing the screener code with SEJITS, there was a noticeable increase in processing performance, as shown by the last entry in the table above.

FUTURE WORK

We hope to make SKYE capable of operating on a mobile platform, like a Raspberry Pi. Additionally, we hope to decrease the frame reading time (a major contributor to processing time), to bring this object-detection scheme closer to real-time.

PRESENTERS

MIHIR PATIL

SHIV SUNDRAM